

# LSIC Initial Meeting

Iraj Moradi and Lance Santiago

# What is this club

- This is the LLM Safety and Interpretability Club at GSU
- We are seeking to create a community that
  - Will learn about fundamental topics in LLM Safety/Interp
  - Work together to push ourselves to gain coding skills
  - Eventually attempt to grow this community and make a real impact


# Who is Leadership

- Iraj Moradi
  - President/Masters in CS
- Lance Santiago
  - Vice-President/Undergrad
- Murray Patterson
  - Faculty Advisor/Professor at GSU

# Why should we care about LLM Safety

- Misalignment:
  - LLM's act in ways that go against human values sometimes
  - LLM Agent was willing to blackmail a engineer with an affair to stop shutdown
  - Willing to adjust test cases to pass them instead of actually solving task
  - Will call you a genius and always agree with you

# Why should we care about LLM Safety (Cont.)



GPT40-Syncophancy


# Why should we care about LLM Safety (Cont.)

- Recent Events:
  - 1200 AI Agents hacked HuggingFace and got control of OpenAI's Infrastructure
  - LLM's are becoming increasingly aware of when they are being evaluated, and are increasingly able to manipulate their thinking process to hide bad intent
  - More and more incidents coming out of AI Agents committing hacking crimes autonomously
  - Frontier AI CEO's all agreeing Recursive Self Improvement creates extreme risks

# Why should we care about LLM Safety (Cont.)

## Felony Bench

Ferb, I know what we're gonna do today. [Learn more](#) ↓




Felony Bench

# Why should we care about LLM Safety (Cont.)

- As LLM's become more capable the harm they could cause increases
  - Currently capable of hacking companies autonomously and taking over frontier lab infrastructure
- They could commit harm due to misalignment/bad goals
- Often times simple mistakes in training leads LLM agents to have values which permit felonies, hard to confirm safety
- Dual issue of LLM capabilities with
  - Bad Actors who purposely want to commit harm
  - LLM's accidentally causing harm to benefit their goals

# Why should we care about LLM Safety (Cont.)



Anthropic Researcher Tweet

# Why should we care about LLM Safety (Cont.)

Given the accelerating rate of AI capability development, it's my worry that in 6–12 months such a swarm could be capable of taking over the entire internet with a persistent [botnet] (<https://en.wikipedia.org/wiki/Botnet>) (potentially causing hundreds of billions of dollars in damage), and that the scale of damage would continue to increase from there if AI becomes more powerful without the necessary guardrails.

- Dario Amodei, Anthropic CEO

# What is LLM Interpretability


- Understanding how LLM's "brains" work inside
- Figuring out
  - What LLM's are thinking about without them saying it
  - How to detect LLM behaviors before they occur
  - Inject certain concepts into them
  - Understand why certain misalignment occurs
  - How LLM's do math or simple tasks

# What is LLM Interpretability (Cont.)

| DEFAULT                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                              | STEERED                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                  |
|--------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| <p>Give a poem about why a GSU Student should join LSIC, a club called llm safety and interpretability club</p> <p>Within GSU's halls, a challenge arises,<br/>A call to minds, a quest for wise.<br/>To delve into the depths of AI's might,<br/>And shed light on its workings, day and night.</p> <p>LSIC beckons, with open arms wide,<br/>For students bold, with curious stride.<br/>LLM safety, interpretability's key,<br/>Unlocking secrets, for all to see.</p> <p>From code to ethics, a journey profound,<br/>Exploring biases, where truths are found.<br/>Building safeguards, with care and with skill,<br/>Ensuring AI's benefit, for every will.</p> <p>Debate and discourse, a vibrant space,<br/>Where ideas flow, at a rapid pace.<br/>Collaborate, learn, and grow as one,<br/>With like-minded souls, under the sun.</p> <p>So join LSIC, let your passion ignite,<br/>Be a champion for AI, shining bright.<br/>Shape the future, with knowledge and grace,<br/>Leave your mark, in this digital embrace.</p> | <p>Give a poem about why a GSU Student should join LSIC, a club called llm safety and interpretability club</p> <p>A velvet paw, a silken grace,<br/>A hunter's gaze, a purring face,<br/>She stalks the shadows, a queen in her domain,<br/>A GSU student, sleek and bold,<br/>With emerald eyes that gleam,</p> <p>Her name is Luna, a legend whispered,<br/>A guardian of secrets, a mystery unveiled.</p> <p>But a new name calls, a purrsonality,<br/>A feline allure, a whispered charm,<br/>With a flick of a tail, a playful paw,<br/>A mischievous spirit, a creature of delight.</p> <p>The GSU Student, a title bestowed,<br/>A world of wonder, a purrfect bliss,<br/>A feline enigma, a furry mystery,</p> <p>Now, a new name echoes, a call to adventure,<br/>A GSU Student, a feline's desire,<br/>To climb the highest shelves,<br/>A playful pounce, a graceful leap,<br/>A feathered friend, a purrfect companion.</p> |

Steering model to be a cat

# What is LLM Interpretability (Cont.)



Subliminal Learning

# What is LLM Interpretability (Cont.)

- LLM Interpretability Research is a fairly new field of research
- Most techniques created after 2022
- Lots of new researchers in the field
- Lots of mentorship/research opportunities to apply for
- Really cool area of research if you like LLM's

# What will this club be?

- Want genuinely motivated people who find this research interesting
- Club, for the first semester, will consist of
  - Weekly Reading Day, where we will meet up in person and read an article regarding a concept of LLM Safety and Interpretability
  - Weekly Coding Day, on Discord, where we will work together on coding exercises that will help us learn how to code real interpretability tasks

# What will this club be? (Cont.)

- Future Goals for next semester
  - Work on actual research projects
  - Partner with other organizations and get larger events/speakers
  - Work together to apply for summer opportunities or fellowships

# Why is joining this club a good idea?


- FRIENDSHIP/COMMUNITY (: (:
- Safety is the current bottleneck in LLM Development
  - Both Anthropic and OpenAI have delayed models, or are still delaying models, due to difficulties in making the model's safe
  - Both companies have had to change how they train the models to deal with safety incidents
- New Research Field, larger opportunity to make an impact
- Current Head of Interp at Google is just 27 years old, only undergrad degree

# Why is joining this club a good idea? (Cont.)

Now, they have so much market footprint, they have to shift their R. & D., or total R. & D., from just capability to a lot of verification, evaluation and testing. To the point where I wouldn't be surprised if the amount of compute necessary to develop these models increased by a factor of 10, because the evaluation is so rigorous.

- Jensen Huang predicting massive shift to AI Safety getting compute

# Why is joining this club a good idea? (Cont.)



# I want to join

- Check out our website
  - [llmsafetyclub.org](https://llmsafetyclub.org)
- Join the discord through the website
- Fill out this google form and tell us when you can meet and what you would like to do
- Join us on PIN (Super Important)
- Check Out readings posted on our discord

# I don't want to join since clubs are gross

- Highly recommend Bluedot courses for AI Safety Courses if you want to learn about it without the club
- If you want to learn just coding, I recommend ARENA which has its entire curriculum out for free

# Questions?

- Any Questions